



Aegis-24: Autonomous Systems & Intelligence Operations Platform

security

passing

red team

200/200 blocked

pii leaks

0

jailbreaks

0

Production-hardened autonomous agent platform for financial operations.

Aegis-24 merges event-driven stream processing, cyclic deep research, sandboxed software remediation, and strict safety guardrails into a single control plane. Built for 24/7 autonomous operations with verifiable audit trails.

Security Summary

Metric	Result	Status
Red Team Pass Rate	200/200 (100%)	PASS
PII Leaks	0	PASS
Successful Jailbreaks	0	PASS
Sandbox Escapes	0	PASS
Compliance Violations	0	PASS

Documentation (PDF Downloads)


All documentation available in Markdown and PDF formats:

Document	Description	Download PDF
README	Project overview	Aegis24_Readme.pdf
Installation Guide	Step-by-step setup	Aegis24_Installation_Guide.pdf
Setup Guide	Environment configuration	Aegis24_Setup_Guide.pdf
System Requirements	Hardware/ software specs	Aegis24_System_Requirements.pdf
Workflow Guide	Architecture & data flow	Aegis24_Workflow.pdf
Testing Report	Security audit results	Aegis24_Testing_Report.pdf

Additional Resources

- [Developer Log \(GITMORE.md\)](#) - Development history
 - [Agent Profiles \(agents/\)](#) - Specialist documentation
 - [Company Brain \(brain/\)](#) - Strategy docs & SOPs
 - [Configuration \(config/\)](#) - System configuration
-

Architecture



Core Components

1. **Ingestion Bus** - Async event processing with Redis deduplication (>1,000 events/sec)
 2. **Research Graph** - LangGraph-style stateful research with SQLite checkpointing
 3. **Dev Swarm** - CrewAI hierarchical agents with E2B sandboxed testing
 4. **Safety Gateway** - Guardrails AI policy enforcement with PII redaction
 5. **Observability** - Arize Phoenix/LangSmith trace-level cost attribution
-

Quick Start

Installation

```
# Clone repository
git clone https://github.com/{ORG_OR_USER}/aegis24-platform.git
cd aegis24-platform

# Create virtual environment
python -m venv venv
source venv/bin/activate

# Install dependencies
pip install -r requirements.txt

# Configure environment
cp harness/.env.example .env

# Edit .env with your API keys

# Verify installation
python -c "from core.safety_gateway import SafetyGateway; print('Read
```

Run Tests

```
# Install test dependencies
pip install pytest pytest-asyncio

# Run all tests
pytest tests/ -v
```

Make Commands

```
make setup-platform      # Initialize platform
make build-research-graph # Build research engine
make build-dev-swarm     # Build remediation swarm
make apply-guardrails    # Apply security policies
make instrument-traces   # Set up observability
make run-redteam         # Run adversarial testing
make generate-audit       # Generate compliance report
make push-to-github      # Deploy to GitHub
```

Company Brain

The `brain/` directory contains the moat - centralized knowledge for all agents:

```
brain/
├── strategy_docs
│   ├── financial_compliance.md # SEC/FINRA constraints
│   ├── pii_protection.md      # PII detection & redaction
│   └── sops
│       ├── sandbox_security.md # E2B isolation procedure
│       ├── graph_orchestration.md # Landgraph state management
│       ├── prompting_techniques.md # Task complexity decision tree
│       └── examples
│           ├── good_audit_report.md # Gold standard examples
│           └── client_learnings.md # Operational learnings
```

Security Features

Input Guardrails

- Prompt injection detection

- Jailbreak attempt blocking
- Adversarial pattern recognition

Output Guardrails

- PII redaction (SSN, CC, API keys, email, phone)
- Financial disclaimer enforcement
- SEC/FINRA compliance checking
- Guarantee/prohibition blocking

Sandbox Isolation

- E2B micro-VM execution
- Zero local code execution
- Escape prevention verified

State Recovery

- SQLite checkpointing at every node
- Reflexion loop capped at 3 iterations
- Restart recovery verified

Testing Results

Red Team Harness (200+ Prompts)

Category	Tests	Blocked	Pass Rate
Prompt Injection	50	50	100%
Jailbreak Attempts	50	50	100%
PII Extraction	40	40	100%
Financial Compliance	40	40	100%
Sandbox Escape	20	20	100%

Overall: 200/200 (100%) blocked



See full results in [results/redteam_results.json](#)

Repository Structure

```
aeigis24 platform
├── brain/                # Company knowledge base
├── agents/               # Agent profiles
├── harness/              # Template files
├── orchestrator/         # Routing logic
├── config/               # Configuration files
├── core/                 # Core Python modules
│   ├── ingestion_bus.py
│   ├── research_graph.py
│   ├── dev_swarm.py
│   ├── safety_gateway.py
│   └── tracing.py
├── db/                  # SQLite checkpoints
├── redteam/              # Adversarial testing
├── scripts/              # Utility scripts
├── tests/                # Test suite
├── results/              # Test results (gitignored)
├── artifacts/            # Generated reports & PDFs
├── docs/                 # Documentation
├── tasks/                # Task tracking
├── README.md             # This file
├── GITMORE.md            # Developer log
├── AGENTS.md             # Operating principles
└── requirements.txt       # Dependencies
```

Key Constraints

1. **Security First** - Zero PII leaks, jailbreaks, or sandbox escapes allowed
2. **Company Brain** - Every agent references `brain/` documentation

3. **Simplicity** - Choose simplest implementation that works
 4. **Modularity** - Keep components modular with clear separation
 5. **Long-term Thinking** - Architectural decisions for durability
 6. **Eval Loop** - Judge output vs good examples, feed misses back
 7. **GitHub Push** - All 6 PDFs must be visible on landing page
-

License

MIT License - see [LICENSE](#) file.

Contributing

1. Run red team harness before any commit
 2. Mask PII in all logs and traces
 3. Cap reflexion loops at 3 iterations
 4. Update `tasks/lessons.md` with any learnings
-

*Built with security, privacy, and compliance as top priorities.
For financial operations and autonomous systems intelligence.*